

# GPGPU Composition with OCaml

Array 2014

Mathias Bourgoïn - Emmanuel Chailloux

June 13, 2014

## OCaml and GPGPU frameworks are very different

### GPGPU frameworks are

- Highly Parallel
- Architecture Sensitive
- Very Low-Level
- Complex to program
- Very efficient


### OCaml is

- Mainly Sequential
- Multi-platform/architecture
- Very High-Level
- Easy to program
- Cannot benefit from parallel architectures

## Idea

- Provide GPGPU programming in their favorite language to OCaml developers.
- Use OCaml to develop high level abstractions for GPGPU.
- Make GPGPU programming safer and easier.

# Overview



## Stream Processing with OCaml



### Features

- Targets Cuda/OpenCL frameworks with OCaml
- Unify these two frameworks
- Abstract memory transfers
- “Lazy” on demand transfers

## What we want

- Simple to express
- Predictable performance
- Easily extensible
- Current high performance libraries
- Optimisable
- Safer

## Two Solutions

### Interoperability with Cuda/OpenCL kernels

- Higher optimisations
- Compatible with current libraries
- Less safe

### A DSL for OCaml : Sarek

- Easy to express
- Easy transformation from OCaml
- Safer

## Sarek Vector Addition

```
let vec_add = kern a b c n ->  
  let open Std in  
  let idx = global_thread_id in  
  if idx < n then  
    c.[<idx>] <- a.[<idx>] + b.[<idx>]
```

## Sarek features

- Monomorphic
- Imperative
- Specific GPGPU globals
- Portable
- ML-like syntax
- Type inference
- Static type checking
- Static compilation to OCaml code
- Dynamic compilation to Cuda and OpenCL
- Exposes its internal representation to the host

# Simple example : vector addition

## SPOC & Sarek

```
open Spoc
let vec_add = kern a b c n ->
  let open Std in
  let idx = global_thread_id in
  if idx < n then
    c.[<idx>] <- a.[<idx>] + b.[<idx>]

let dev = Devices.init ()
let n = 1_000_000
let v1 = Vector.create Vector.float64 n
let v2 = Vector.create Vector.float64 n
let v3 = Vector.create Vector.float64 n

let block = {blockX = 1024; blockY = 1; blockZ = 1}
let grid={gridX=(n+1024-1)/1024; gridY=1; gridZ=1}


let main () =
  random_fill v1;
  random_fill v2;
  Kirc.gen vec_add;
  Kirc.run vec_add (v1, v2, v3, n) (block,grid) dev.(0);
  for i = 0 to Vector.length v3 - 1 do
    Printf.printf "res[%d] = %f; " i v3.[<i>]
  done;
```

OCaml  
No explicit transfers  
Type inference  
Static typing  
Portable  
Heterogeneous

# Sequential composition : real-world example

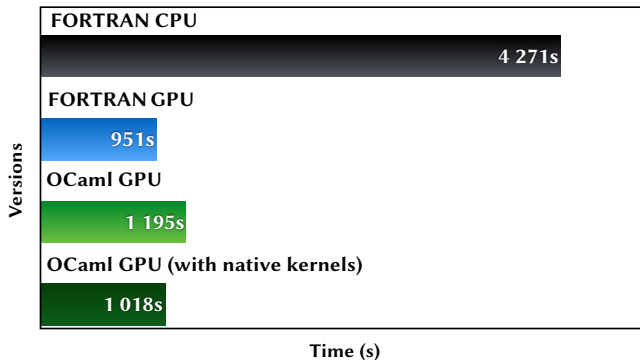
## PROP

- Included in the 2DRMP suite
- HPC prize for Machine Utilization, awarded by the UK Research Councils' HEC Strategy Committee, 2006
- Simulates  $e^-$  scattering in H-like ions at intermediate energies
- PROP Propagates a  $\mathcal{R}$ -matrix in a two-electrons space
- Computations mainly implies matrix multiplications
- Computed matrices grow during computation
- Programmed in Fortran
- Compatible with sequential architectures, HPC clusters, super-computers





# Results: PROP



SPOC+Sarek achieves 80% of hand-tuned Fortran performance.  
SPOC+external kernels is on par with Fortran (93%)

Type-safe  
Memory manager + GC

30% code reduction  
No more transfers

# Functional composition

## Why?

- Helps describe complex algorithms
- Composition constructs can provide better optimizations

## Problem

Kernels are procedures

Difficult to identify input/output automatically (with external native kernels)

## Two solutions

- Analyze and transform Sarek internal representation
- Provide skeletons to associate kernels with their inputs/outputs

# Solution 1 : Sarek transformations

## Using Sarek

Transformations are OCaml functions transforming Sarek AST :

Example:

```
map (kern a → b)
```

Scalar computations ( $'a \rightarrow 'b$ ) are transformed into vector ones ( $'a \text{ vector} \rightarrow 'b \text{ vector}$ ).

## Vector addition

```
let v1 = Vector.create Vector.float64 10_000
and v2 = Vector.create Vector.float64 10_000 in
let v3 = map2 (kern a b → a + b) v1 v2

val map2 :
  ('a → 'b → 'c) sarek_kernel →
  ?dev:Spoc.Devices.device →
  'a Spoc.Vector.vector →
  'b Spoc.Vector.vector → 'c Spoc.Vector.vector
```

# Solution 1 : Sarek transformations

```
sort (kern a b → a - b) vec1  
val sort : ('a → 'a → int) sarek_kernel → 'a vector → unit
```

## Injection into sort kernel

```
let bitonic_sort = kern v j k →  
  let open Std in  
  let i = thread_idx_x +  
    block_dim_x*block_idx_x in  
  let ixj = Math.xor i j in  
  let mutable temp = 0. in  
  if ixj >= i then (  
    if (Math.logical_and i k) = 0 then (  
      if v.[<i>] - v.[<ixj>] > 0 then  
        (temp := v.[<ixj>];  
         v.[<ixj>] <- v.[<i>];  
         v.[<i>] <- temp)  
      else if v.[<i>] - v.[<ixj>] <= 0 then  
        (temp := v.[<ixj>];  
         v.[<ixj>] <- v.[<i>];  
         v.[<i>] <- temp);)
```

```
while !k <= size do  
  j := !k lsr 1;  
  while !j > 0 do  
    run bitonic_sort  
      (vec1,!j,!k)  
      device;  
    j := !j lsr 1;  
  done;  
  k := !k lsl 1 ;  
done;
```

— Host composition

# Solution 2 : Parallel skeletons

## A skeleton combines

- a kernel
- an execution environment
- an input
- an output

## Two running functions:

- *run* : runs on one device
- *par\_run* : tries running on a list of devices

- Explicitly describes relations between kernels/data
- Provides automatic optimizations

# Examples

## Skeleton

```
(* 'a : environment, 'b : input, 'c : output *)  
val MAP : 'a external_kernel -> 'b vector -> 'c vector -> ('a, 'b, 'c) skeleton  
val run : ('a, 'b, 'c) skeleton -> 'a -> 'c vector
```

- Automatic grid/block mapping on GPU
- Automatic parallelization on multiple GPUs

## Composition

```
val PIPE : ('a, 'b, 'c) skeleton -> ('d, 'c, 'e) skeleton -> ('f, 'b, 'e) skeleton
```

- Automatic overlapping of transfers by computations

## Our solution

- High-level multiparadigm programming language : **OCaml**
- Runtime library with vectors and lazy transfers : **SPOC**
- Easily extensible dedicated DSL : **Sarek**
- Kernel transformations combining kernels and host composition
- Functional composition and automatic optimizations *via* parallel skeletons

## Portable approach

Requires a high level programming language with :

- Multiple paradigms
- Customizable garbage collector
- C interoperability

## Improve composability

- Use Sarek to provide deeper transformations
- Build more skeletons
- Add a cost model to Sarek
- Target highly heterogeneous systems

## High performance web-client programming

- Using `js_of_ocaml`
- Targeting WebCL
- Demands to translate SPOC's low-level C code to javascript
- Helps develop rich multimedia applications with intensive computations
- Eases accessibility : perfect playground for GPGPU/HPC courses



# Thanks



Emmanuel Chailloux  
Jean-Luc Lamotte

Open-Source distribution : <http://www.algo-prog.info/spoc/>  
Or install it via [OPAM](#), the OCaml Package Manager  
SPOC is compatible with x86\_64: Unix (Linux, Mac OS X), Windows

For more information  
[mathias.bourgoin@lip6.fr](mailto:mathias.bourgoin@lip6.fr)



# A Little Example



CPU RAM



GPU0 RAM



GPU1 RAM

## Example

```
let dev = Devices.init ()
let n = 1_000_000
let v1 = Vector.create Vector.float64 n
let v2 = Vector.create Vector.float64 n
let v3 = Vector.create Vector.float64 n

let k = vec_add (v1, v2, v3, n)
let block = {blockX = 1024; blockY = 1; blockZ = 1}
let grid={gridX=(n+1024-1)/1024; gridY=1; gridZ=1}

let main () =
  random_fill v1;
  random_fill v2;
  Kernel.run k (block,grid) dev.(0);
  for i = 0 to Vector.length v3 - 1 do
    Printf.printf "res[%d] = %f; " i v3.[<i>]
  done;
```

# A Little Example



v1  
v2  
v3  
CPU RAM



GPU0 RAM



GPU1 RAM

## Example

```
let dev = Devices.init ()
let n = 1_000_000
let v1 = Vector.create Vector.float64 n
let v2 = Vector.create Vector.float64 n
let v3 = Vector.create Vector.float64 n

let k = vec_add (v1, v2, v3, n)
let block = {blockX = 1024; blockY = 1; blockZ = 1}
let grid={gridX=(n+1024-1)/1024; gridY=1; gridZ=1}

let main () =
  random_fill v1;
  random_fill v2;
  Kernel.run k (block,grid) dev.(0);
  for i = 0 to Vector.length v3 - 1 do
    Printf.printf "res[%d] = %f; " i v3.[<i>]
  done;
```

# A Little Example



v1  
v2  
v3  
CPU RAM



GPU0 RAM



GPU1 RAM

## Example

```
let dev = Devices.init ()
let n = 1_000_000
let v1 = Vector.create Vector.float64 n
let v2 = Vector.create Vector.float64 n
let v3 = Vector.create Vector.float64 n

let k = vec_add (v1, v2, v3, n)
let block = {blockX = 1024; blockY = 1; blockZ = 1}
let grid={gridX=(n+1024-1)/1024; gridY=1; gridZ=1}

let main () =
  random_fill v1;
  random_fill v2;
  Kernel.run k (block,grid) dev.(0);
  for i = 0 to Vector.length v3 - 1 do
    Printf.printf "res[%d] = %f; " i v3.[<i>]
  done;
```

# A Little Example



v1  
v2  
v3  
CPU RAM



GPU0 RAM



GPU1 RAM

## Example

```
let dev = Devices.init ()
let n = 1_000_000
let v1 = Vector.create Vector.float64 n
let v2 = Vector.create Vector.float64 n
let v3 = Vector.create Vector.float64 n

let k = vec_add (v1, v2, v3, n)
let block = {blockX = 1024; blockY = 1; blockZ = 1}
let grid={gridX=(n+1024-1)/1024; gridY=1; gridZ=1}

let main () =
  random_fill v1;
  random_fill v2;
  Kernel.run k (block,grid) dev.(0);
  for i = 0 to Vector.length v3 - 1 do
    Printf.printf "res[%d] = %f; " i v3.[<i>]
  done;
```

# A Little Example



CPU RAM



GPU0 RAM

v1  
v2  
v3



GPU1 RAM

## Example

```
let dev = Devices.init ()
let n = 1_000_000
let v1 = Vector.create Vector.float64 n
let v2 = Vector.create Vector.float64 n
let v3 = Vector.create Vector.float64 n

let k = vec_add (v1, v2, v3, n)
let block = {blockX = 1024; blockY = 1; blockZ = 1}
let grid={gridX=(n+1024-1)/1024; gridY=1; gridZ=1}

let main () =
  random_fill v1;
  random_fill v2;
  Kernel.run k (block,grid) dev.(0);
  for i = 0 to Vector.length v3 - 1 do
    Printf.printf "res[%d] = %f; " i v3.[<i>]
  done;
```

# A Little Example



v3  
CPU RAM



v1  
v2  
GPU0 RAM



GPU1 RAM

## Example

```
let dev = Devices.init ()
let n = 1_000_000
let v1 = Vector.create Vector.float64 n
let v2 = Vector.create Vector.float64 n
let v3 = Vector.create Vector.float64 n

let k = vec_add (v1, v2, v3, n)
let block = {blockX = 1024; blockY = 1; blockZ = 1}
let grid={gridX=(n+1024-1)/1024; gridY=1; gridZ=1}

let main () =
  random_fill v1;
  random_fill v2;
  Kernel.run k (block,grid) dev.(0);
  for i = 0 to Vector.length v3 - 1 do
    Printf.printf "res[%d] = %f; " i v3.[<i>]
  done;
```

# Example

## Power iteration

### SPOC

```
while (iter<IterMax)&&(max_n>eps) do
  let x=A*x0 in
  let m = max(x) in
  let x=u/m in
  let n = abs(x - x0) in
  max_n <- max(n);
  x0<-x; iter<-iter+1;
done
```

### Skeletons

```
while (iter<IterMax)&&(max_n>eps) do
  let x= map ( * x0) A in
  let m = reduce (max) x in
  let x= map ( / m) u in
  let n = map (abs) x-x0 in
  max_n <- reduce max n;
  x0<-x; iter<-iter+1;
done
```

## Composition

```
while (iter<IterMax)&&(max_n > eps) do
  let m= pipe (map ( *x0)) (reduce max) A in
  max_n <- pipe
    (pipe
      (map ( / m)
        (map (abs(- x0)))))
    (reduce max) u;
  x0<-x; iter<-iter+1;
done
```



# Example

## Power Iteration

### SPOC

```
while (iter<IterMax)&&(max_n>eps) do
  let x=A*x0 in
  let m = max(x) in
  let x=u/m in
  let n = abs(x - x0) in
  max_n <- max(n);
  x0<-x; iter<-iter+1;
done
```

### Skeletons

```
while (iter<IterMax)&&(max_n>eps) do
  let x= map ( * x0) A in
  let m = reduce (max) x in
  let x= map ( / m) u in
  let n = map (abs) x-x0 in
  max_n <- reduce max n;
  x0<-x; iter<-iter+1;
done
```

## Composition

```
while (iter<IterMax)&&(max_n > eps) do
  let m= pipe (map ( *x0)) (reduce max) A in
  max_n <- pipe
    (pipe
      (map ( / m)
        (map (abs(- x0)))))
    (reduce max) u;
  x0<-x; iter<-iter+1;
done
```

# Example

## Power Iteration

### SPOC

```
while (iter<IterMax)&&(max_n>eps) do
  let x=A*x0 in
  let m = max(x) in
  let x=u/m in
  let n = abs(x - x0) in
  max_n <- max(n);
  x0<-x; iter<-iter+1;
done
```

### Skeletons

```
while (iter<IterMax)&&(max_n>eps) do
  let x= map ( * x0) A in
  let m = reduce (max) x in
  let x= map ( / m) u in
  let n = map (abs) x-x0 in
  max_n <- reduce max n;
  x0<-x; iter<-iter+1;
done
```

## Composition

```
while (iter<IterMax)&&(max_n > eps) do
  let m= pipe (map ( *x0)) (reduce max) A in
  max_n <- pipe
    (pipe
      (map ( / m)
        (map (abs(- x0)))))
    (reduce max) u;
  x0<-x; iter<-iter+1;
done
```

# Composition benchmarks

